



Start For Free



DEC 4, 2023

READ TIME: 4 MIN

Introducing Data Portal in Stream Governance

Technology > Confluent



Olivia Greene

Senior Product Marketing Manager

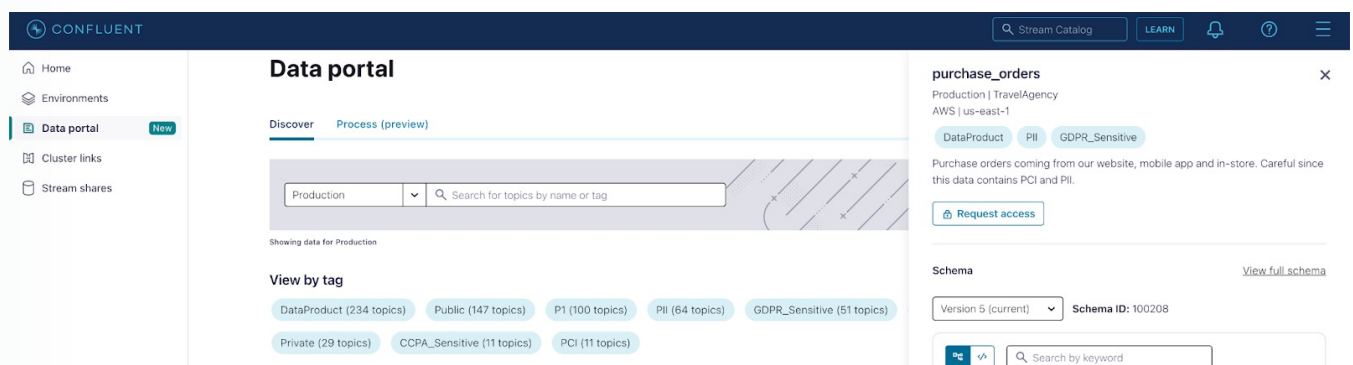


David Araujo

Director of Product Management



Today, we're excited to announce the general availability of [Data Portal](#) on Confluent Cloud. Data Portal is built on top of [Stream Governance](#), the industry's only fully managed data governance suite for Apache Kafka® and data streaming. The developer-friendly, self-service UI provides an easy and curated way to find, understand, and enrich all of your data streams, enabling users across your organization to build and launch streaming applications faster.



By clicking "Accept All Cookies", you agree to the storing of cookies on your device to enhance site navigation, analyze site usage, and assist in our marketing efforts. [Cookie Notice](#)

Accept All Cookies

Reject All

Cookies Settings

Data Portal is a self-service interface for discovering, exploring, and accessing Kafka topics on Confluent Cloud.

Building streaming applications with open source Kafka can be slow and inefficient when there's a lack of visibility into what data exists, where it comes from, and who can grant access. Data Portal leverages the capabilities of [Stream Catalog](#) and [Stream Lineage](#) to empower data users to interact with their organization's Kafka data streams safely, efficiently, and collaboratively.

With Data Portal, you can:

- **Search and discover existing topics** across the organization with the help of topic metadata and get a drill-down view of the data they hold.
- **Seamlessly and securely request access to topics** through an approval workflow that connects the data user with the data owner, who can approve the request.
- **Set up clients and query data with Apache Flink®** to enrich your topics and build new streaming applications and pipelines.

Ready to get started? If you already use Confluent Cloud, you can access Data Portal simply by [logging in](#) to your account. You must have a [Stream Governance package enabled](#) for the cloud environments you want displayed in Data Portal. Check out the [quick start guide](#) to see how you can get Stream Governance up and running in just a few clicks.

If you're not yet using Confluent Cloud, you can [try Data Portal for free](#) by creating an account and setting up your first cluster.

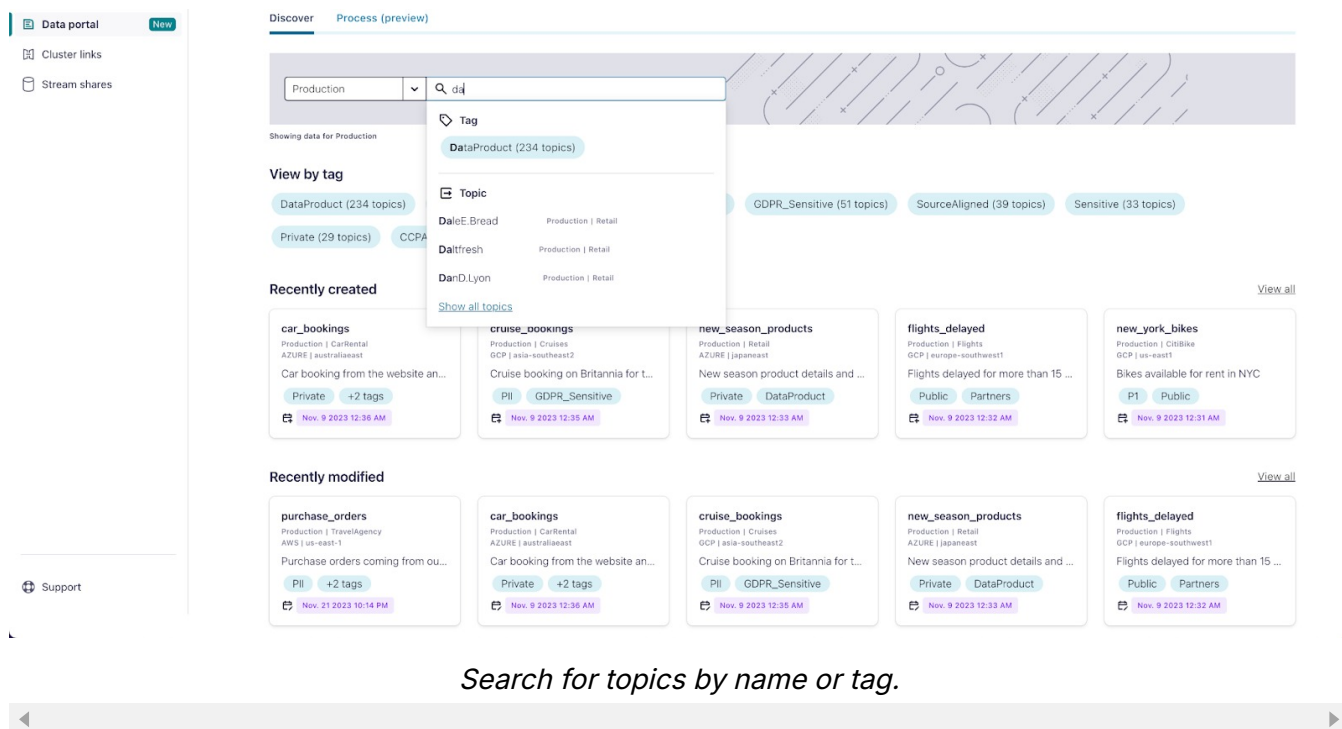
To learn more, join us for the [Stream Governance webinar](#) with a technical demo that showcases the full capabilities of Data Portal and Stream Governance on Confluent Cloud.

[Join Us For A Demo](#)

Let's dive into how to get up and running with Data Portal.

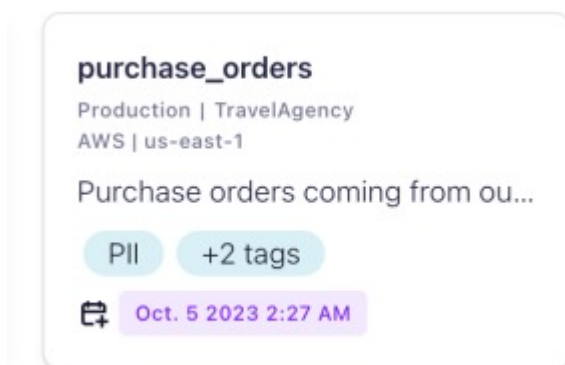
Search and discover existing topics

When you log in to Confluent Cloud, the **Data Portal** tab will reflect a unified view of all available Kafka topics by environment. You can search for topics by name or [tag](#), or browse topics by tag, creation date, and modified date.



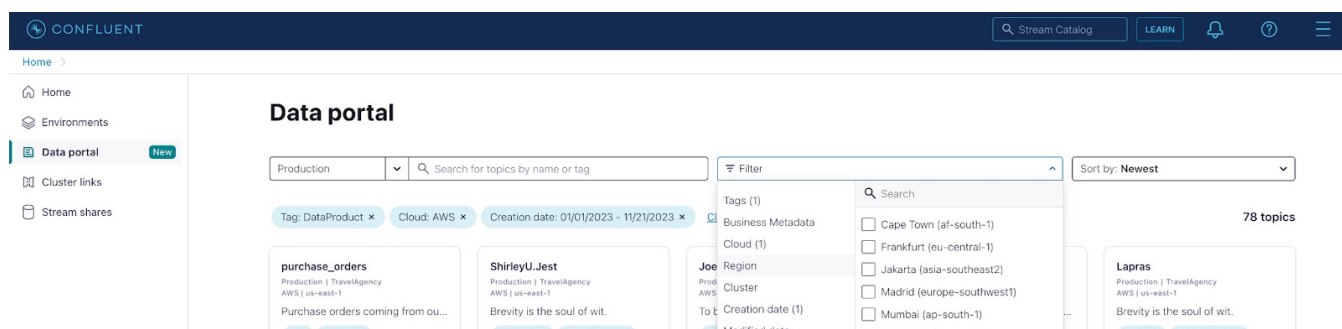
Search for topics by name or tag.

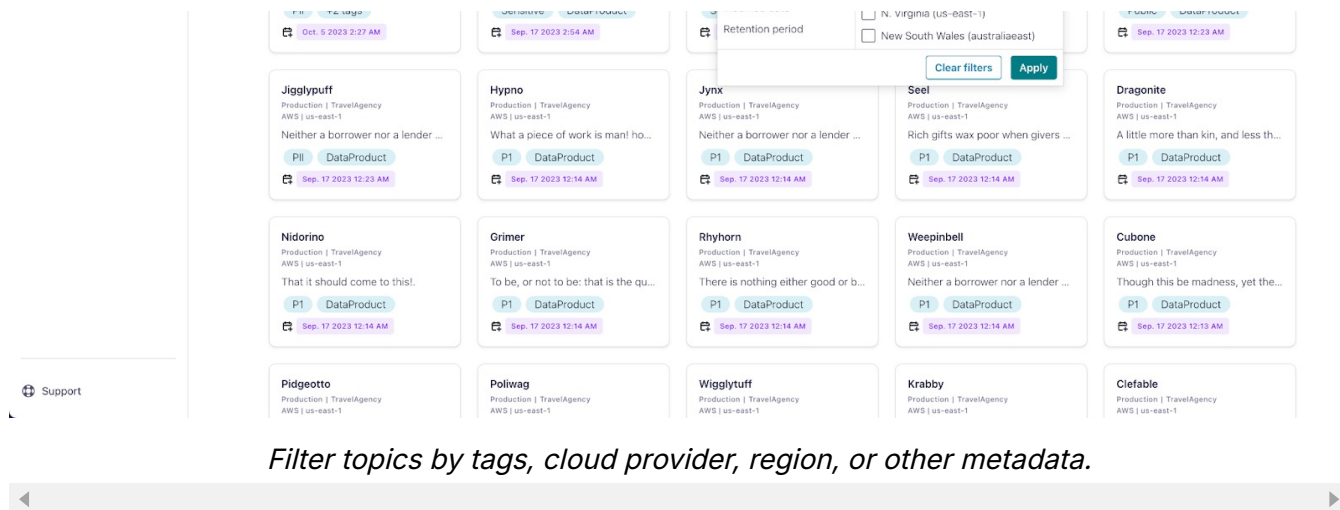
Each topic card summarizes the topic with the name, data location (environment, cluster, cloud provider, and region), description, tags, and when it was created or modified.



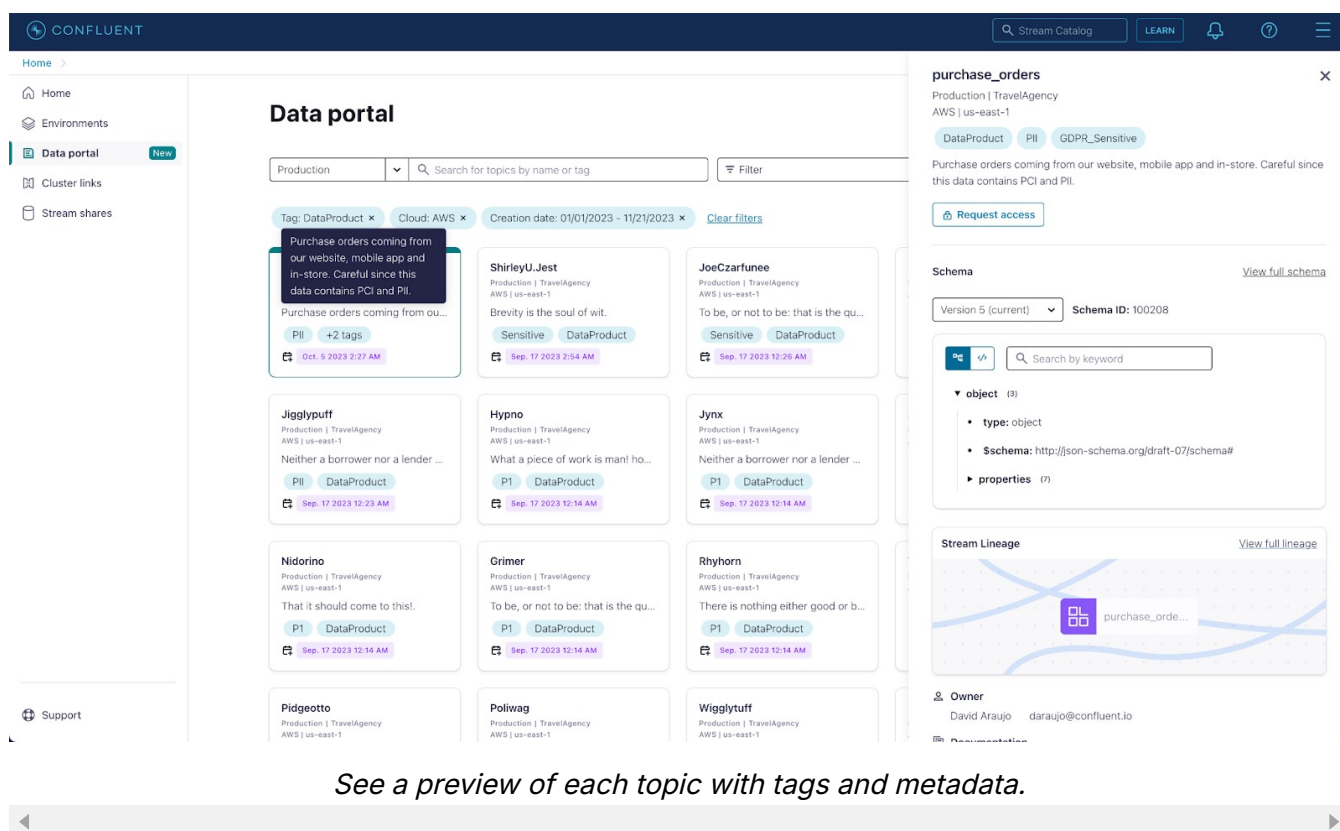
See a preview of each topic with tags and metadata.

However, this comprehensive view may be overwhelming if you're sifting through hundreds of topics. Data Portal's search feature allows you to hone in on the most relevant data for your use case or project. Search for topics by name or tag, or filter by tags, **business metadata**, cloud provider, region, and other metadata.

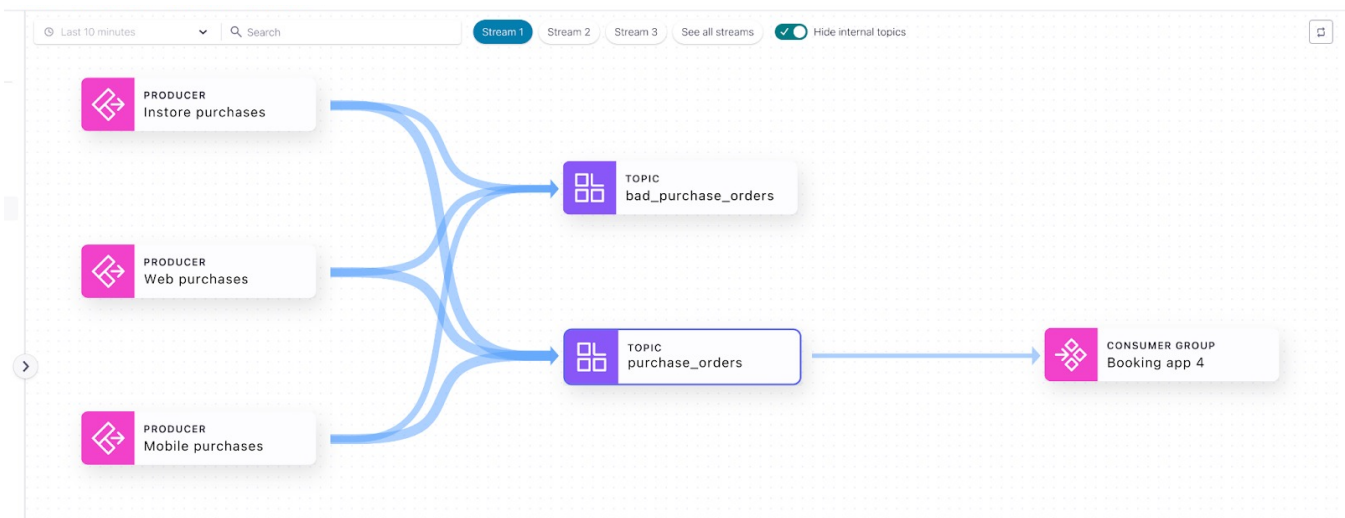




Once you've identified a topic that piques your interest, click on the card to learn more about it. Clicking on the card reveals a side panel with additional metadata. In the top section of this panel, you'll see the location of the topic, its tags, and a description of the data it stores. Below, you can see the schema with its fields (without viewing the actual data), where you can view the structure of the data stored in the topic. You'll also find a link to the lineage of the topic, information about the owner of the topic, business metadata appended to the topic, and finally, the technical metadata of the topic (created date, retention period, etc).



When you click on the **Stream Lineage** section, you get a complete, **end-to-end data flow visualization** of the upstream and downstream components from the topic.



Expand the Stream Lineage graph for an end-to-end view of the data streams.

Seamlessly request access to topics

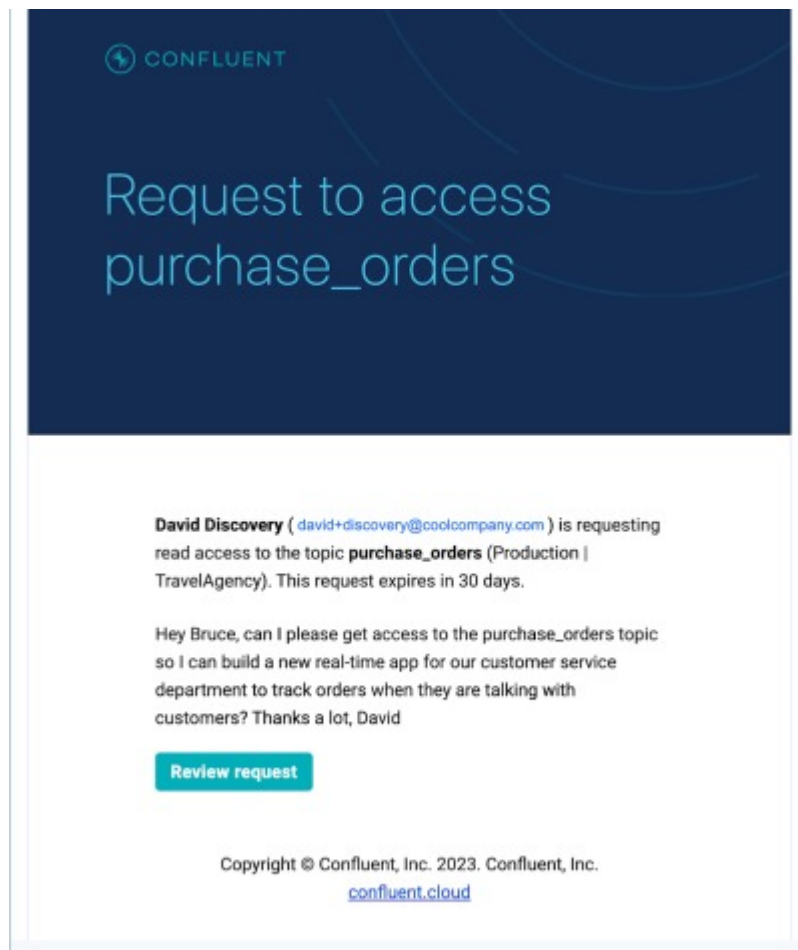
So you've identified the topic(s) you want to access. Now what? Instead of spending valuable cycles pinging various colleagues to find out who can grant access to the data you need, you can use the Data Portal to request access to the topic directly from the Confluent Cloud UI.

Clicking **Request access** on the topic side panel triggers an approval workflow that connects the user with the data owner via email (if a [topic owner email was set on the topic metadata](#)). Select the permissions you require and (optionally) leave a message for the approver describing your request.

The screenshot shows a modal dialog titled 'Request access to purchase_orders'. Below the title, it says 'Select access and send request to the topic owner'. There are two radio button options: 'Read only' (which is selected) and 'Read and Write'. Below these is a text area labeled 'Message (optional)'. The message area contains the text: 'Hey Bruce, can I please get access to the purchase_orders topic so I can build a new real-time app for our customer service department to track orders when they are talking with customers? Thanks a lot, David'. At the bottom right of the dialog are two buttons: 'Cancel' and 'Request'.

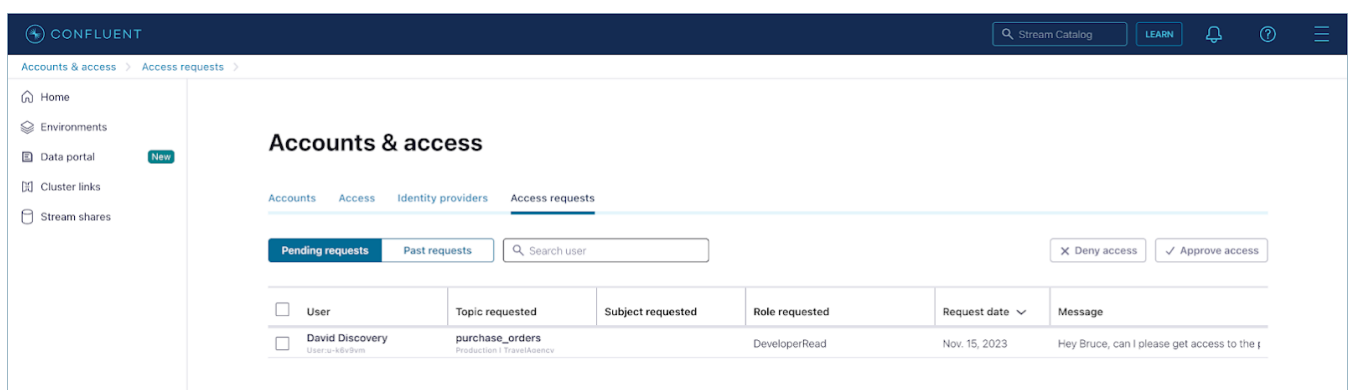
Request access to topics with a self-service approval workflow within the Confluent Cloud UI.

Once the request is submitted, the topic owner will receive an email to review the request.



Email notifications let the topic owner know when a request is submitted.

When the topic owner clicks **Review request** in the email, they are redirected to the Confluent Cloud **Access requests** UI.



Data owners and admins can view requests and manage access to topics in a single view.

This simple workflow provides important self-service capabilities for both data owners and data seekers. Streamlining access request management eliminates the need for administrators to manually manage manual permission assignments, while at the same time ensuring that security and access controls remain in place – a welcome sight as your users and topics scale to the tens and hundreds (and beyond)!

Set up clients and query data with Flink

Let's assume the data owner granted you access to the topic – now when you click on the topic card, you can view the last message produced and set up a client or query the topic with [Flink SQL](#) directly from the UI.

The screenshot shows the Confluent Data portal interface. On the left is a sidebar with navigation links: Home, Environments, Data portal (selected), Cluster links, and Stream shares. The main area is titled 'Data portal' and contains a search bar and filters. Below the filters is a grid of topic cards. Each card displays the topic name, production environment, a brief description, and tags like 'PII' and '+2 tags'. The 'purchase_orders' topic is highlighted. To the right of the grid, a detailed view of the 'purchase_orders' topic is shown, including a 'Message' section with a JSON payload.

Data portal

Production | TravelAgency
AWS | us-east-1

Tags: DataProduct x Cloud: AWS x Creation date: 01/01/2023 - 11/21/2023 x Clear filters

purchase_orders
Production | TravelAgency
AWS | us-east-1
Purchase orders coming from ou...
PII +2 tags
Oct 5 2023 2:27 AM

ShirleyU.Jest
Production | TravelAgency
AWS | us-east-1
Brevity is the soul of wit.
Sensitive DataProduct
Sep 17 2023 2:54 AM

JoeCzarfunee
Production | TravelAgency
AWS | us-east-1
To be, or not to be: that is the qu...
Sensitive DataProduct
Sep 17 2023 12:26 AM

Jigglypuff
Production | TravelAgency
AWS | us-east-1
Neither a borrower nor a lender ...
PII DataProduct
Sep 17 2023 12:23 AM

Hypno
Production | TravelAgency
AWS | us-east-1
What a piece of work is man! ho...
P1 DataProduct
Sep 17 2023 12:14 AM

Jynx
Production | TravelAgency
AWS | us-east-1
Neither a borrower nor a lender ...
P1 DataProduct
Sep 17 2023 12:14 AM

Nidorino
Production | TravelAgency
AWS | us-east-1
That it should come to this!
P1 DataProduct
Sep 17 2023 12:14 AM

Grimer
Production | TravelAgency
AWS | us-east-1
To be, or not to be: that is the qu...
P1 DataProduct
Sep 17 2023 12:14 AM

Rhyhorn
Production | TravelAgency
AWS | us-east-1
There is nothing either good or b...
P1 DataProduct
Sep 17 2023 12:14 AM

Pidgeotto
Production | TravelAgency
AWS | us-east-1

Poliwag
Production | TravelAgency
AWS | us-east-1

Wigglytuff
Production | TravelAgency
AWS | us-east-1

purchase_orders View topic
Production | TravelAgency
AWS | us-east-1
DataProduct PII GDPR_Sensitive
Purchase orders coming from our website, mobile app and in-store. Careful since this data contains PCI and PII.
Set up client Query

Message View all messages

Value	Header	Key
1	{	
2	"contact": {	
3	"email": "curt.conroy@yahoo.com",	
4	"mailingAddress": "160 Luetgen Corner, South	
5	Groverport, NJ 08706",	
6	"phoneNumber": "917-722-6746"	
7	},	
8	"fare": {	
9	"ticketPrice": 66.91	
10	},	
11	"flight": {	
12	"arrivalAirport": "VEAB",	
13	"departureAirport": "NFFA",	
14	"departureDateTime": 1699697719482,	
15	"flightNumber": "Su-15"	
16	},	
17	"itinerary": {	
18	"totalTravelTime": "6"	
19	},	
20	"passenger": {	
21	"dateOfBirth": "Mon Sep 25 17:05:06 UTC 1972",	
22	"firstName": "Jess",	
23	"lastName": "Strosin"	
24	},	
25	"payment": {	
26	"creditCardInfo":	
27	"21EguMgnFk97K80guAQILMIU7arInSr55d2ffxP2PFTdRvn	
	Gfmo2ce",	
	"paymentDate": 1699618128265	
	},	

View messages, set up a client, and query data with Flink.

Clicking **Set up client** will redirect you to the **Clients** page in the UI, which will walk you through [setting up an application](#) in your programming language of choice.

Clicking **Query** takes you to the new [Flink SQL workspace](#) with a topic query ready to run.

The screenshot shows the Confluent Flink SQL workspace interface. The top bar includes navigation links: Home, Environments, Production (selected), and a workspace ID: workspace-2023-11-16-003823. The main area displays a SQL query: `SELECT * FROM 'Production'.TravelAgency.'purchase_orders' LIMIT 10;`. Below the query, there are buttons for 'Stop' and 'Run'. The results section shows a table with 10 rows of data.

workspace-2023-11-16-003823

David Araujo AWS | US-EAST-1
User Account AWS-east1-pool2 Running

Use catalog Use database

```
1 SELECT * FROM 'Production'.TravelAgency.'purchase_orders' LIMIT 10;
```

START TIME: 2023-11-16T00:44:45Z STATEMENT STATUS: Stopped STATEMENT ID: workspace-2023-11-16-003823-1 RESULTS: 10

KEY	SEAT	FARE	FLIGHT	PASSENGER	CONTACT	ITINERARY	PAYME
	(23V)	(81.47)	(WAAB, , ULAA, Ka-25)	(Marisol, Torp, Tue Sep 11 ...	((795) 494-8784, 93374 ...	(8)	(V1xSN
	(8h)	(20.25)	(ZBBB, , WA44, Li-2)	(Heriberto, Strosin, Sat Fe...	(1-936-521-1004, Suite 8...	(5)	(Oc2nN
	(23m)	(9.71)	(FAMO, , VEBT, SSJ)	(Cary, Rowe, Tue Dec 18 1...	(1-889-800-2433, 6082 ...	(8)	(Oz5Clr
	(3c)	(96.04)	(UUOO, , EBLG, Tu-2)	(Karl, Ryan, Sat Jul 15 05:...	(906-994-2645, 446 Eug...	(8)	(lbyPmY
	(9f)	(76.41)	(CYBA, , WAAB, Boeing 7...	(Fabian, Stroman, Fri Dec ...	((476) 859-1795, Suite 09...	(3)	(kd7BN

Check out the latest blog post on [our serverless Flink service](#) to learn more about how you can effortlessly filter, join, and enrich your Kafka data streams in-flight.

What's next?

We often hear how critical Stream Governance is to our customers' data in motion journey, and with Data Portal, we're excited to bring an enhanced user experience to the product. We look forward to expanding our Stream Governance suite further and adding exciting new features in quarters to come.

Check out **Data Portal** in [Confluent Cloud](#), today. If you haven't already, sign up for a [free trial of Confluent Cloud](#) and create your first cluster to explore new topics and create streaming pipelines and applications.

Interested in learning more? Be sure to register for the upcoming [Stream Governance webinar](#), where we'll share a technical demo that showcases the full capabilities of Data Portal and Stream Governance on Confluent Cloud.

[Register For Stream Governance Webinar](#)



Olivia is a product marketer at Confluent, focused on driving the adoption of Confluent Cloud and Confluent Platform and empowering businesses to harness the full potential of data streaming. Before Confluent, she led programs and go-to-market initiatives across infrastructure, storage, and edge technologies at Red Hat.



David Araujo is the Director of Product Management for Stream Governance at Confluent. As an engineer turned product manager, over the years he has worked across multiple industries and continents, mostly in the data management and strategy space. He holds a master's and bachelor's degree in computer science from the University of Evora in Portugal.

How to Build a Data Mesh in Confluent Cloud with Stream

Governance

In this webinar, you'll learn how to build a data mesh and ensure high data quality, discoverability, and compatibility for your real-time Apache Kafka® data streams on Confluent Cloud using the latest Stream Governance capabilities.

[Register Now >](#)

Free Course: Get Started with Stream Governance

Get an in-depth introduction to streaming data governance in our course on Confluent Developer!

[Start the Course >](#)


Did you like this blog post? Share it now



[Technology](#) < [Confluent](#)

Subscribe to the Confluent blog

Subscribe



Confluent Unlocks the Full Power of Event Streams with Stream Governance

SEP 14, 2021

Data governance initiatives aim to manage the availability, integrity, and security of data used across an organization. With the explosion in volume, variety, and velocity of data powering the modern [...]

DAVID ARAUJO



How to Find, Share, and Organize Your Data Streams with Stream Catalog

JUL 14, 2022

In a modern data stack, being able to discover, understand, organize and reuse data is key to obtaining the most value from your data. These requirements have led to the [...]

ROBERT YOKOTA

DAVID ARAUJO



[Terms & Conditions](#) | [Privacy Policy](#) | [Do Not Sell My Information](#) | [Modern Slavery Policy](#) | [Cookie Settings](#)

Copyright © Confluent, Inc. 2014-2024. Apache®, Apache Kafka®, Kafka®, Apache Flink®, Flink®, and associated open source project names are trademarks of the Apache Software Foundation